

Figure 9: The Gradient Bottleneck — Architecture, Not Scale, Is the Key
All Models Encode Bias (probe $\approx 0.68+$); Only Llama Gradients Reveal It

