

Figure 15: Generalization to Safety Domain (BeaverTails)
Influence Functions Detect Safety Bias Beyond RLHF Preference Data

