

Figure 13: Top-k Suspicious Sample Detection
Enrichment over Random Baseline — Llama-3.2-3B Achieves $3.7\times$ at $k=50$

